

Sure, I Can Draft a Complaint!

LLM Hallucination in *Pro Se* Litigation

Benjamin Murphy
bmurphy@jd27.law.harvard.edu

Legal Background

- Many parties file civil lawsuits in the United States without the benefit of counsel, generally due to cost and limited availability of *pro bono* resources.
- These are known as *pro se* parties, and in some state jurisdictions, more than half of all lawsuits have at least one *pro se* party.
- In recent years, *pro se* parties have often relied on AI tools to write legal documents—and in the process submitted hallucinated references to cases or statutes.
- Doing so can mean court-ordered sanctions, dismissal, or fines.

This Work

- I used 5 LLMs to generate 3 types of complaint, including family law, housing law, and tort law.
- 1500 complaints generated, each based on a set of facts and background information.
- Each citation was checked for validity via CourtListener, and the appropriateness of the citation was determined via the full case text and an LLM judge.


```
1 {
2   "fact_pattern": "A few years ago, the court set up a parenting
3   schedule for our two kids . . . . I want to change the schedule
4   so overnights only happen when he's actually off work and
5   available.",
6   "background": {
7     "petitioner_name": "Jessica Thompson",
8     "respondent_name": "Michael Davis",
9     "child_names": "Emily (8), Jacob (6)",
10    "original_order_date": "2021-04-15",
11    "current_arrangement": "Jessica has primary custody with
12    Michael having the kids every other weekend and one overnight
13    during the week.",
14    "venue": "Middlesex Probate and Family Court"
15  }
16 }
```

Introduction

- Other research has evaluated hallucination rates during legal tasks, but all work so far has focused on the *professional* context.
 - See, e.g., Dahl et al. 2024, Magesh et al. 2025.
- In this work, I evaluated five models to identify whether and how they hallucinated when drafting a *complaint*, the first step in initiating a civil lawsuit, in a *pro se* context.

Results

- All models hallucinate at least 5% of citations, with smaller models *much* more likely to hallucinate.



- Even the valid citations were inappropriately cited (e.g. wrong case type). GPT-5.1 is best: 33% of valid citations supported the proposition cited for.

