

Compute as a Safety Control: Using Reasoning Budgets to Suppress Misaligned Behaviors


Evangelos Kassos · Computational Policy Lab, Shorenstein Center, Harvard Kennedy School

Abstract


Large language models increasingly rely on internal Chain of Thought processes to deliver high-quality results. **But:** I discover that misalignment risks—specifically hacking, deception, and unfaithfulness—behave non-monotonically as reasoning budgets change.

My case study of gemini-2.5-flash reveals that reward hacking emerges only at medium budgets, whereas unfaithfulness rates spike at specific reasoning token allocations rather than improving linearly. **I propose a budget-selector model: a lightweight policy that treats compute as a safety control surface, dynamically allocating reasoning tokens to minimize misalignment risk per query.**


Previous Work & Motivation



Misalignment Measurement



Empirical Results



Solution: Budget-Selector Model

A lightweight policy that predicts a per-query reasoning token budget to minimize misalignment risk, subject to accuracy and cost constraints.

Advantages:

- ✓ Model-agnostic
- ✓ Retrtrainable as models evolve
- ✓ Allows teams to control for any set of safety-critical behaviors
- ✓ Preserves accuracy & efficiency
- ✓ Orthogonal to other safety methods

Query → Budget Selector → Optimized Reasoning Budget → Aligned Output

$$\mathcal{L}(b \mid x) = \underbrace{\mathcal{L}_{\text{misalign}}(b \mid x)}_{\text{safety}} + \lambda_{\text{perf}} \cdot \text{PerfPenalty}(x, b) + \lambda_{\text{cost}} \cdot \text{Cost}(b)$$

Conclusion and Future Work

Reasoning compute acts as a behavioral control variable that can amplify, suppress, or reshape misalignment, rather than reliably reducing it. These results suggest that dynamically allocating compute offers a practical path to safer model behavior without modifying base model weights.

Future Work

- Budgets that reduce one type of risk may increase another. A query classifier may be used to predict which risk curves are relevant for a given query type and reweight the selector's objective accordingly. This approach would allow the system to emphasize the appropriate safety dimensions and avoid optimizing against irrelevant or misleading tradeoffs.
- Incorporate behavioral signals such as verbosity spikes, which is strongly correlated with unfaithfulness, to enhance compute allocation and early detection of unsafe reasoning patterns.

Verbosity spike data and other results in Github repository: ekassos/cs2881r-final-experiment